

# Convolutional Neural Network (Recap)

Naeemullah Khan

[naeemullah.khan@kaust.edu.sa](mailto:naeemullah.khan@kaust.edu.sa)

Prashant Aparajeya



جامعة الملك عبد الله  
للعلوم والتقنية  
King Abdullah University of  
Science and Technology

KAUST Academy  
King Abdullah University of Science and Technology

May 20, 2025

1. Definition
2. Convolutional Layer
3. Activation Functions
4. Pooling Layers
5. Padding & Strides
6. Normalization (Batch Norm)
7. Regularization (Dropout)
8. Flattening & Fully Connected Layers
9. Loss Functions & Optimizers
10. Architectures
  - 10.1 LeNet
  - 10.2 AlexNet

10.3 VGG

10.4 InceptionNet

10.5 ResNet

10.6 MobileNet

11. Transfer Learning

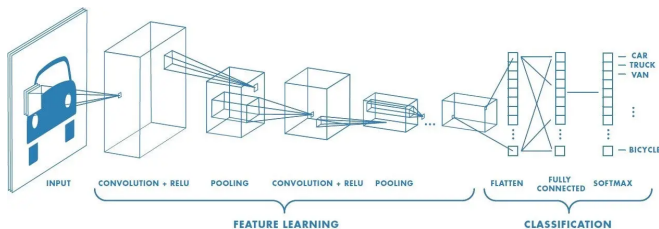
12. Evaluation Metrics

13. Further Reading

- ▶ Convolutional Neural Networks (CNNs) are a class of deep learning models specifically designed for processing structured grid data, such as images.
- ▶ They are particularly effective for tasks like image classification, object detection, and segmentation.
- ▶ CNNs leverage the spatial structure of images by using convolutional layers to automatically learn hierarchical features.
- ▶ The architecture typically consists of convolutional layers, activation functions, pooling layers, and fully connected layers.
- ▶ CNNs are known for their ability to capture local patterns and translate them into higher-level representations.

## What is a CNN?

A CNN is a deep network of neurons with learnable filters that perform convolution operations on inputs, usually images, to extract hierarchical features.

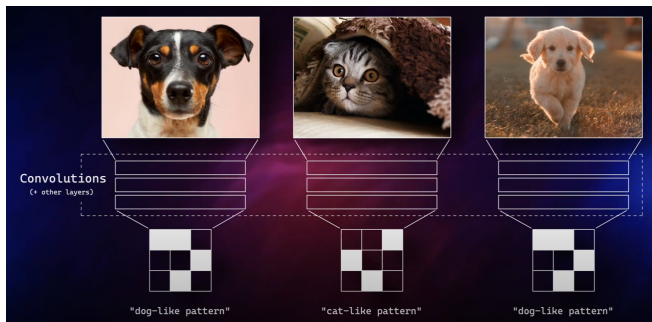


## Why use CNNs?

Parameter sharing and sparse connectivity reduce number of parameters and improve spatial feature extraction.


## What makes a Convolutional Neural Network?

Characterised by “Convolutional Layer” – they are able to detect “abstract features” and “almost ideas within the image”




# Components of a CNN


## Convolution Layers




## Pooling Layers



## Fully-Connected Layers



## Activation Function



## Normalization

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \varepsilon}}$$

## Operation:

- ▶ Element-wise multiply filter with image patch and sum  $\rightarrow$  feature map.

## Hyperparameters:

- ▶ kernel size
- ▶ number of filters
- ▶ stride
- ▶ padding

### Listing 1: Code snippet (PyTorch)

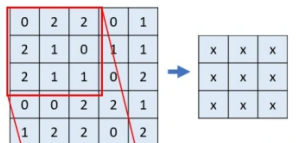
```
import torch.nn as nn

conv = nn.Conv2d(in_channels=3, out_channels=16,
                 kernel_size=3, stride=1, padding=1)

output = conv(input_tensor) # input_tensor: [
    batch_size, 3, H, W]
```



# CNN - Convolutional Layer



Without padding, the edges of the image are only partially processed, and the result of convolution is smaller than the original image size

$$\text{Filter} = F \times F$$


0	0	0
0	0	1
1	1	0

1
---

bias

width =  $W \times W$

padding =  $P$



stride =  $S$

$$\text{Filter} = F \times F$$

0	0	0
0	0	1
1	1	0

1
---

bias

1. Convolution result size =  $(W - F + 2P) / S + 1$
2.  $(W - F + 2P) / S + 1$  should be an integer
3. If you set  $S = 1$ , then setting  $P = (F - 1) / 2$  will generate convolution result size equal to the image size.

# Quick Exercise (5 mins)

Let's find out what this can give us:

- ▶ Padding = 0
- ▶ Stride = 1



Note: Once you traverse entire image/matrix it will give you a matrix calls Feature Map or Activation Map.